

# On-device Deep Learning

Zhaojun Ni


School of Information Science and Technology



# Key Reference

- Han Song's Group @ MIT
- Yoo Hoi-Jun's Group @ KAIST
- Shi Yuanming's Group @ ShanghaiTech

# Deep Learning Going “Tiny”




**Cloud AI**

Data centers



**Mobile AI**

Smartphones



**Tiny AI**

IoT Devices/Microcontrollers

Memory  
(Activation)

32GB

4GB

320kB

Storage  
(Weights)

~TB/PB

512GB

1MB



# On-device Deep Learning is Essential

- AI systems need to adapt to new sensor data for **customization** and **continual learning**
- To protect user's **private** data or achieve higher performance in **real-world** applications
- However, tiny devices contain only **limited** computation capability



# Outline: From Design to Inference

## □(Design) Efficient Neural Architecture

- Neural Architecture Search(NAS)
- Resource-constrained model specialization

## □(Training) Efficient Training Hardware

- Dataflow optimizations
- External memory access reduction
- Computation optimizations

## □(Inference) Efficient Inference System

- Memory-efficient inference engine
- Edge training and inference system

# #1. Neural Architecture Search(NAS)

- not only reduce FLOPs or latency
- extend the **search space** to fit the tiny **resource constraints**

$$S' = \text{kernel size} \times \text{expansion ratio} \times \text{depth} \times \text{input resolution } \underline{R} \times \text{width multiplier } \underline{W}$$

- Different  $R$  and  $W$  for different hardware capacity


$$R=260, W=1.4$$



$$R=224, W=1.0$$



$$R=?, W=?$$




F412/F743/H746/...  
256kB/320kB/512kB/...



# #1. Neural Architecture Search(NAS)

- Analyzing FLOPs distribution of satisfying models in each search space:


**Larger FLOPs -> Larger model capacity -> More likely to give higher accuracy**



# #2. Resource-constrained model specialization


## Kernel Size

Smaller kernel takes centered weights



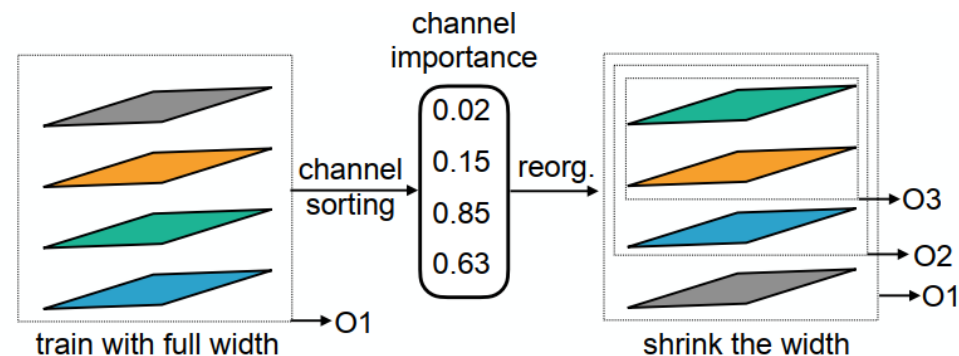
## Depth

Allow later layers in each unit to be skipped to reduce the depth



## Width

Keep the most important channels when shrinking via channel sorting



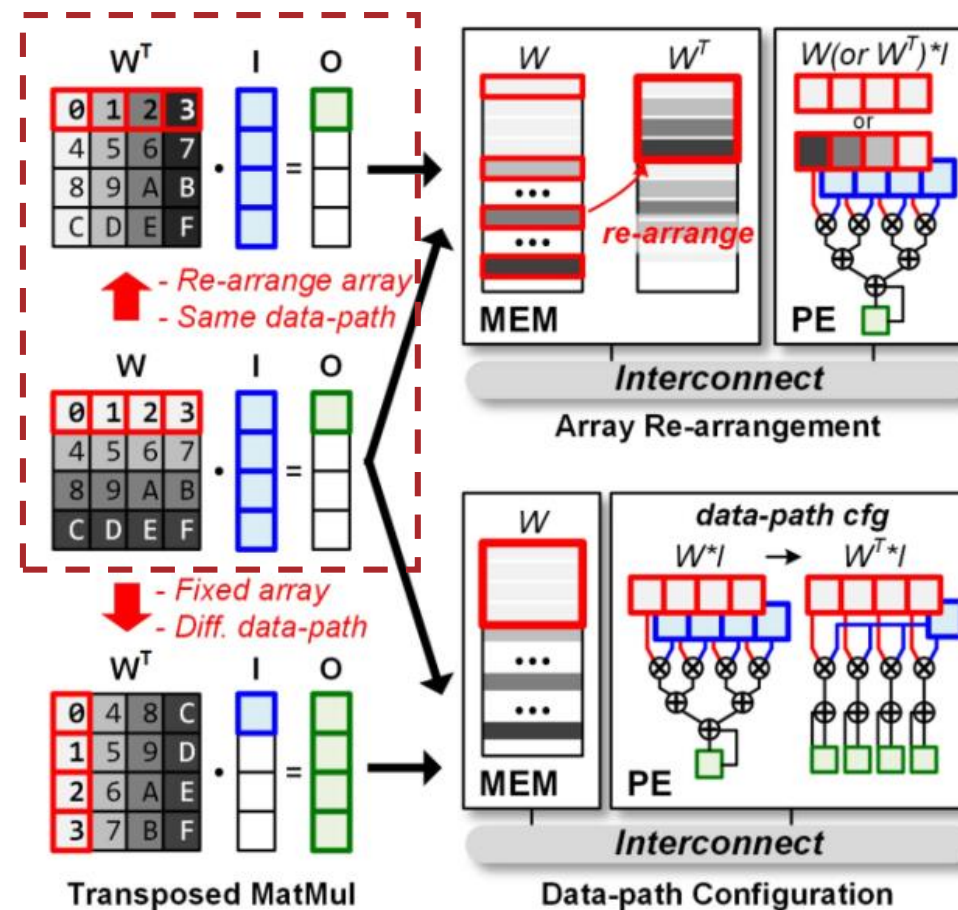


# #3. Dataflow optimizations

## Data array re-arrangement

Modify the **memory** layout

- loads tiled weights from external memory and stores to on-chip SRAM or regfiles in transposed order
- read array for both transposed and non-transposed manner with a bit-rotator
- custom on-chip SRAM which can read data as transposed

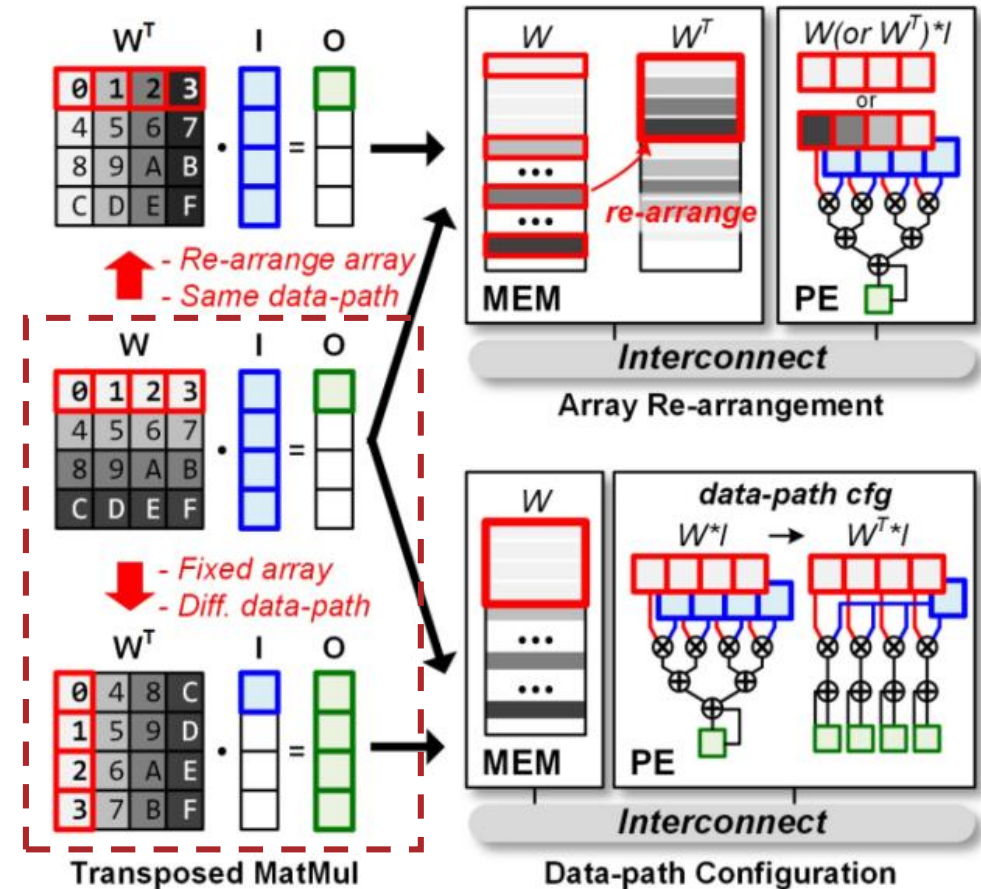


# #3. Dataflow optimizations

## Data-path configuration

Modify data-path of the **PE array**

- configurable and multiple data-paths in a PE array
- transposed and non-transposed MatMul without redundant memory access by exchangeable feeding paths of input and weight
- different dataflow of DNN training steps with heterogeneous architectures



# #4. External memory access reduction

## Sparse compression

Represent the tensors into a non-zero vector

- Zero-Value Compression (ZVC)
- Compressed Sparse Row (CSR)
- Run Length Encoding (RLE)

Sparse compression

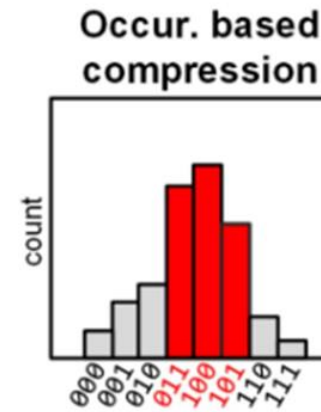
0	5	0	0
1	0	0	0
0	0	0	0
2	0	4	0

	compressed vector
data	[5 1 2 4]
cidx	[1 0 0 2]
ridx	[0 1 3 3]

## Probability of occurrence based compression

Encode more likely to occur values into low bit-width code

- Huffman coding
- Replaced most-frequent exponent values with low-bit code






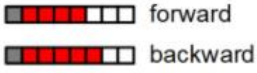
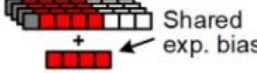

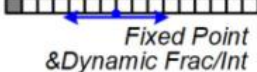
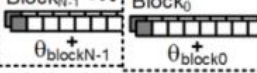
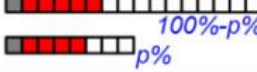
data	p(%)	code
011	35	01
100	32	10
101	31	11
000	0.5	00000
001	1.1	00001
...	...	...



# #4. External memory access reduction

## Reduced Bit-width

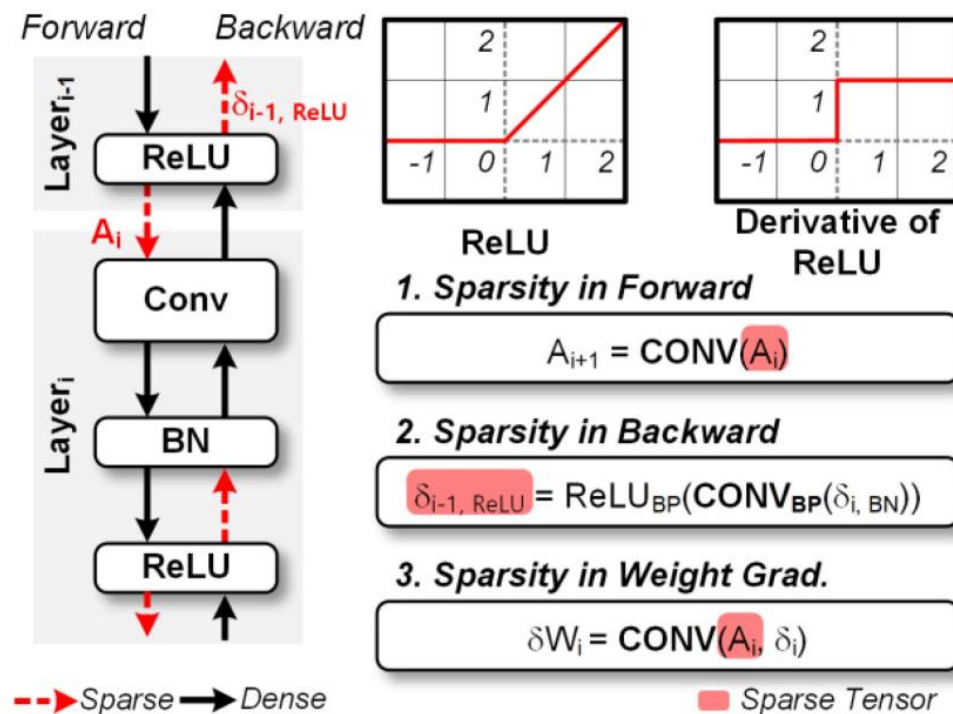
- custom reduced width data-type
- some hardware uses fp16 or fp8
- use low bit-width for representing the stored data
- 1 bit-width for ReLU-Pool layer pairs

Data-type	Opt. Level	Format
Mixed Precision [75]	Network	
bfloat [65]	Network	 Extended exponent
DLfloat [66]	Network	 Extended exponent
Hybrid fp8 [67]	Network	 forward backward
fp8-SEB [59]	Layer	 Shared exp. bias
Flexpoint [76]	Layer	 Shared exp. fxp
SDFXP [62]	Layer	 Fixed Point & Dynamic Frac/Int
LDQ [77]	Neuron	 Block <sub>N-1</sub> ... Block <sub>0</sub> $\theta_{blockN-1}$ ... $\theta_{block0}$
FGMP fp8/fp16 [47]	Neuron	 100%-p% p%

## #5. Computation optimization method

# Sparsity

- replaces fp16 operations with zero input MAC operations and compensates with fp8 operations
- zero output prediction and speculative skipping
- determine the weight to be pruned in the training phase and skip all related operations
- neglect the partial sum of small values



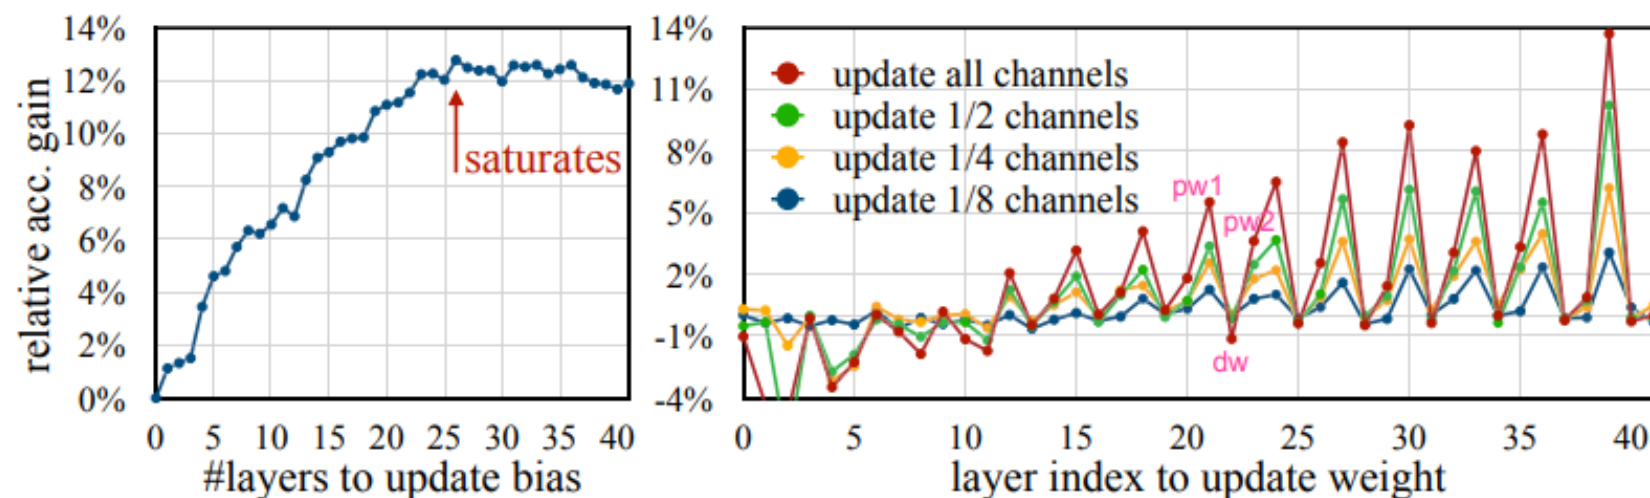


# #5. Computation optimizations

## Sparse layer/tensor update


Prune the gradient during backpropagation and update the model sparsely

Find the right sparse update scheme



(a) Contribution of last  $k$  biases  $\Delta acc_{b_{[k]}}$


(b) Contribution of a certain weight  $\Delta acc_{W_{i,r}}$




(a) full update



(b) bias-only update



(c) sparse layer update



(d) sparse tensor update


# #5. Computation optimizations

## Quantization


- custom bit-precisions
- configurable PE design for various bit-width

- Quantization-Aware Training(QAT)
- Quantization-Aware Scaling(QAS)


$$\tilde{\mathbf{G}}_{\tilde{\mathbf{W}}} = \mathbf{G}_{\tilde{\mathbf{W}}} \cdot s_{\tilde{\mathbf{W}}}^{-2}, \quad \tilde{\mathbf{G}}_{\tilde{\mathbf{b}}} = \mathbf{G}_{\tilde{\mathbf{b}}} \cdot s_{\tilde{\mathbf{W}}}^{-2} \cdot s_{\tilde{\mathbf{x}}}^{-2} = \mathbf{G}_{\tilde{\mathbf{b}}} \cdot s^{-2}$$



(a) Real quantized graph (ours)



(b) Fake quantized graph (QAT)



# #6. Memory-efficient inference engine

## Separated compilation & runtime

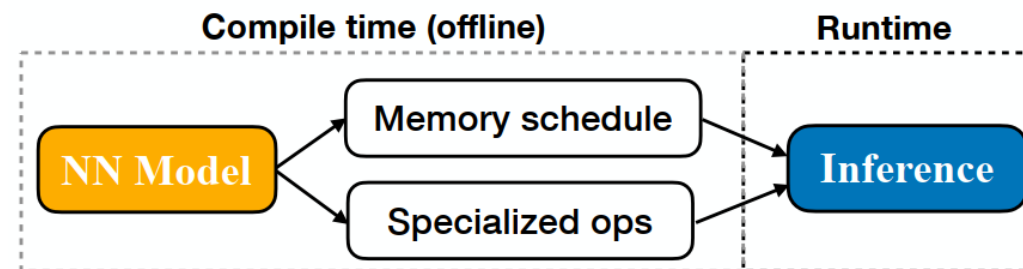
- Reduce overhead

## Code generator-based compilation

- Eliminate overheads of runtime interpretation

## Model-adaptive memory scheduling

- Increase data reuse for each layer



(a) Model-level memory scheduling

$$M = \max (\text{kernel size}_{L_i}^2 \cdot \text{in channels}_{L_i}; \forall L_i \in \mathbf{L})$$

(b) Tile size configuration for Im2col

$$\text{tiling size of feature map width}_{L_j} = \lfloor M / (\text{kernel size}_{L_j}^2 \cdot \text{in channels}_{L_j}) \rfloor$$



# #6. Memory-Efficient inference engine

## Patch-based inference

- Reduce peak memory
- Reduce the receptive field of the patch-based initial stage
- Increase the receptive field of the later stage

## Graph-level optimization

- Minimize memory footprint
- Optimize the overall computation

## Re-order and in-place update

- Gradient updates are immediately applied once calculated
- Intermediate buffers can be released

# #7. Edge training and inference system

## Data partition based edge training systems

- Data is massively distributed over a number of edge devices, and each edge device has only a subset of the whole dataset
- During training, each edge device holds a replica of the complete AI model to compute a local update



(a) Distributed mode



(b) Decentralized mode

# #7. Edge training and inference system

## Model partition based edge training systems

- Each node holds part of the model parameters with small storage size
- Accomplish the model training task or the inference task collaboratively
- Data privacy at each node belongs to different parties
- Heavy communication overhead between edge devices

# #7. Edge training and inference system


## Computation offloading based edge inference systems

- Offload the entire inference task to an edge server


Edge devices should upload their input data to edge servers for inference

- Offload only a part of the task to the edge server

Edge server computes the inference result based on the intermediate value computed by the edge device



(a) Server-based edge inference

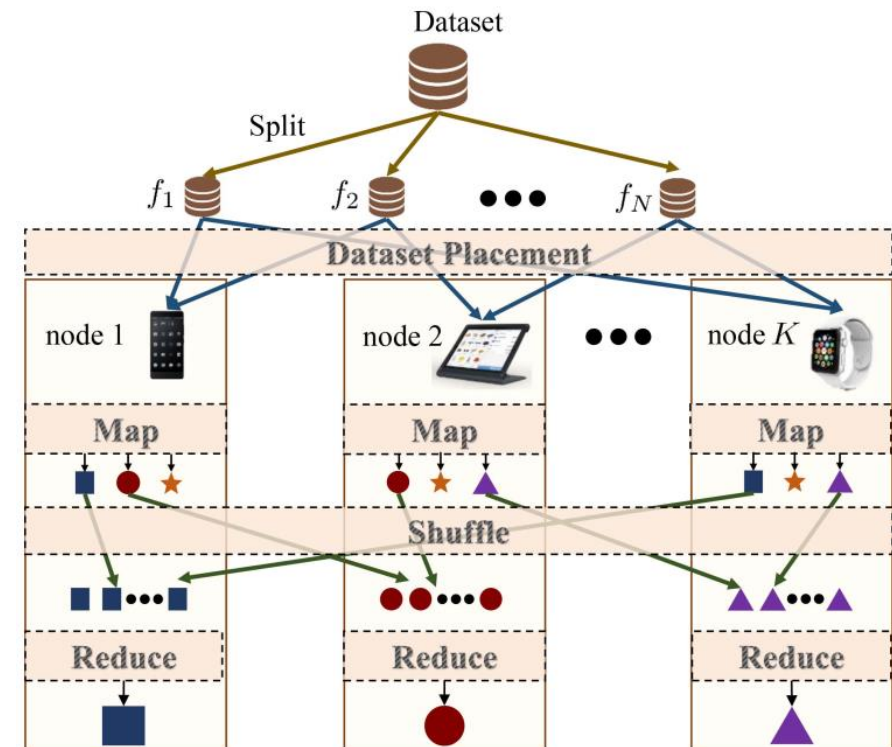


(b) Device-edge joint inference

# #7. Edge training and inference system

## General edge computing systems

- MapReduce
- In the map phase, every computing node computes a map function of the assigned data
- In the shuffle phase, nodes communicate with each other to obtain some intermediate values
- In the reduce phase, each node computes the assigned output function



# THANKS

